

# Analysis of Algorithms

## Linear-Time Sorting Algorithms

1

11/30/2002

## Sorting So Far

- I Insertion sort:
  - n Easy to code
  - n Fast on small inputs (less than ~50 elements)
  - n Fast on nearly-sorted inputs
  - n  $O(n^2)$  worst case
  - n  $O(n^2)$  average (equally-likely inputs) case
  - n  $O(n^2)$  reverse-sorted case

2

11/30/2002

## Sorting So Far

- I Merge sort:
  - n Divide-and-conquer:
    - u Split array in half
    - u Recursively sort subarrays
    - u Linear-time merge step
  - n  $O(n \lg n)$  worst case
  - n Doesn't sort in place

3

11/30/2002

## Sorting So Far

- I Heap sort:
  - n Uses the very useful heap data structure
    - u Complete binary tree
    - u Heap property: parent key > children's keys
  - n  $O(n \lg n)$  worst case
  - n Sorts in place
  - n Fair amount of shuffling memory around

4

11/30/2002

## Sorting So Far

- I Quick sort:
  - n Divide-and-conquer:
    - u Partition array into two subarrays, recursively sort
    - u All of first subarray < all of second subarray
    - u No merge step needed!
  - n  $O(n \lg n)$  average case
  - n Fast in practice
  - n  $O(n^2)$  worst case
    - u Naïve implementation: worst case on sorted input
    - u Address this with randomized quicksort

5

11/30/2002

## How Fast Can We Sort?

- I We will provide a lower bound, then beat it
  - n *How do you suppose we'll beat it?*
- I First, an observation: all of the sorting algorithms so far are *comparison sorts*
  - n The only operation used to gain ordering information about a sequence is the pairwise comparison of two elements
  - n Theorem: all comparison sorts are  $\Omega(n \lg n)$ 
    - u A comparison sort must do  $O(n)$  comparisons (*why?*)
    - u What about the gap between  $O(n)$  and  $O(n \lg n)$

6

11/30/2002

## Decision Trees

- I *Decision trees* provide an abstraction of comparison sorts
  - n A decision tree represents the comparisons made by a comparison sort. Every thing else ignored
  - n (Draw examples on board)
- I *What do the leaves represent?*
- I *How many leaves must there be?*

7

11/30/2002

## Decision Trees

- I Decision trees can model comparison sorts. For a given algorithm:
  - n One tree for each  $n$
  - n Tree paths are all possible execution traces
  - n *What's the longest path in a decision tree for insertion sort? For merge sort?*
- I *What is the asymptotic height of any decision tree for sorting  $n$  elements?*
- I Answer:  $\Omega(n \lg n)$  (now let's prove it...)

8

11/30/2002

## Lower Bound For Comparison Sorting

- Thm: Any decision tree that sorts  $n$  elements has height  $\Omega(n \lg n)$
- What's the minimum # of leaves?*
- What's the maximum # of leaves of a binary tree of height  $h$ ?*
- Clearly the minimum # of leaves is less than or equal to the maximum # of leaves

9

11/30/2002

## Lower Bound For Comparison Sorting

- So we have...
- $n! \leq 2^h$
- Taking logarithms:  
 $\lg(n!) \leq h$
- Stirling's approximation tells us:

$$n! > \left(\frac{n}{e}\right)^n$$

- Thus:  $h \geq \lg\left(\frac{n}{e}\right)^n$

10

11/30/2002

## Lower Bound For Comparison Sorting

- So we have
- $h \geq \lg\left(\frac{n}{e}\right)^n$
- $= n \lg n - n \lg e$
- $= \Omega(n \lg n)$
- Thus the minimum height of a decision tree is  $\Omega(n \lg n)$

11

11/30/2002

## Lower Bound For Comparison Sorts

- Thus the time to comparison sort  $n$  elements is  $\Omega(n \lg n)$
- Corollary: Heapsort and Mergesort are asymptotically optimal comparison sorts
- But the name of this lecture is "Sorting in linear time"!
- How can we do better than  $\Omega(n \lg n)$ ?*

12

11/30/2002

## Sorting In Linear Time

- Counting sort
  - No comparisons between elements!
  - But...* depends on assumption about the numbers being sorted
    - We assume numbers are in the range  $1..k$
  - The algorithm:
    - Input:  $A[1..n]$ , where  $A[j] \in \{1, 2, 3, \dots, k\}$
    - Output:  $B[1..n]$ , sorted (notice: not sorting in place)
    - Also: Array  $C[1..k]$  for auxiliary storage

13

11/30/2002

## Counting Sort

```

1  CountingSort(A, B, k)
2    for i=1 to k
3      C[i] = 0;
4    for j=1 to n
5      C[A[j]] += 1;
6    for i=2 to k
7      C[i] = C[i] + C[i-1];
8    for j=n downto 1
9      B[C[A[j]]] = A[j];
10   C[A[j]] -= 1;
```

Work through example:  $A=\{4\ 1\ 3\ 4\ 3\}$ ,  $k=4$

14

11/30/2002

## Counting Sort

```

1  CountingSort(A, B, k)
2    for i=1 to k
3      C[i] = 0;
4    for j=1 to n
5      C[A[j]] += 1;
6    for i=2 to k
7      C[i] = C[i] + C[i-1];
8    for j=n downto 1
9      B[C[A[j]]] = A[j];
10   C[A[j]] -= 1;
```

What will be the running time?

15

11/30/2002

## Counting Sort

- Total time:  $O(n + k)$ 
  - Usually,  $k = O(n)$
  - Thus counting sort runs in  $O(n)$  time
- But sorting is  $\Omega(n \lg n)$ !
  - No contradiction--this is not a comparison sort (in fact, there are *no* comparisons at all!)
  - Notice that this algorithm is *stable*

16

11/30/2002

## Counting Sort

- | Cool! *Why don't we always use counting sort?*
- | Because it depends on range  $k$  of elements
- | *Could we use counting sort to sort 32 bit integers? Why or why not?*
- | Answer: no,  $k$  too large ( $2^{32} = 4,294,967,296$ )

17

11/30/2002

## Counting Sort

- | *How did IBM get rich originally?*
- | Answer: punched card readers for census tabulation in early 1900's.
  - In particular, a *card sorter* that could sort cards into different bins
    - Each column can be punched in 12 places
    - Decimal digits use 10 places
  - Problem: only one column can be sorted on at a time

18

11/30/2002

## Radix Sort

- | Intuitively, you might sort on the most significant digit, then the second msd, etc.
- | Problem: lots of intermediate piles of cards (read: scratch arrays) to keep track of
- | Key idea: sort the *least* significant digit first

```
RadixSort(A, d)
  for i=1 to d
    StableSort(A) on digit i
```

  - Example: Fig 9.3

19

11/30/2002

## Radix Sort

- | *Can we prove it will work?*
- | Sketch of an inductive argument (induction on the number of passes):
  - Assume lower-order digits  $\{j: j < i\}$  are sorted
  - Show that sorting next digit  $i$  leaves array correctly sorted
    - If two digits at position  $i$  are different, ordering numbers by that digit is correct (lower-order digits irrelevant)
    - If they are the same, numbers are already sorted on the lower-order digits. Since we use a stable sort, the numbers stay in the right order

20

11/30/2002

## Radix Sort

- | *What sort will we use to sort on digits?*
- | Counting sort is obvious choice:
  - Sort  $n$  numbers on digits that range from  $1..k$
  - Time:  $O(n + k)$
- | Each pass over  $n$  numbers with  $d$  digits takes time  $O(n+k)$ , so total time  $O(dn+dk)$ 
  - When  $d$  is constant and  $k=O(n)$ , takes  $O(n)$  time
- | *How many bits in a computer word?*

21

11/30/2002

## Radix Sort

- | Problem: sort 1 million 64-bit numbers
  - Treat as four-digit radix  $2^{16}$  numbers
  - Can sort in just four passes with radix sort!
- | Compares well with typical  $O(n \lg n)$  comparison sort
  - Requires approx  $\lg n = 20$  operations per number being sorted
- | *So why would we ever use anything but radix sort?*

22

11/30/2002

## Radix Sort

- | In general, radix sort based on counting sort is
  - Fast
  - Asymptotically fast (i.e.,  $O(n)$ )
  - Simple to code
  - A good choice
- | To think about: *Can radix sort be used on floating-point numbers?*

23

11/30/2002

## The End

24

11/30/2002